

```
In [1]: import warnings                                     #Importing few important python li
warnings.filterwarnings('ignore')
import pandas as pd
import numpy as np
from numpy import linalg as LA
from sklearn.model_selection import train_test_split
import math
import matplotlib.pyplot as plt
import copy
import math
%matplotlib inline
```

```
In [2]: df = pd.read_csv('heartdataset.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

| | male | age | education | currentSmoker | cigsPerDay | BPMeds | prevalentStroke | prevalentHyp | diabe |
|---|------|-----|-----------|---------------|------------|--------|-----------------|--------------|-------|
| 0 | 1 | 39 | 4.0 | 0 | 0.0 | 0.0 | 0 | 0 | |
| 1 | 0 | 46 | 2.0 | 0 | 0.0 | 0.0 | 0 | 0 | |
| 2 | 1 | 48 | 1.0 | 1 | 20.0 | 0.0 | 0 | 0 | |
| 3 | 0 | 61 | 3.0 | 1 | 30.0 | 0.0 | 0 | 1 | |
| 4 | 0 | 46 | 3.0 | 1 | 23.0 | 0.0 | 0 | 0 | |

```
In [4]: df.shape
```

```
Out[4]: (4238, 16)
```

```
In [5]: df.isnull().sum()
```

```
Out[5]: male                0
age                0
education          105
currentSmoker      0
cigsPerDay         29
BPMeds             53
prevalentStroke    0
prevalentHyp       0
diabetes           0
totChol            50
sysBP              0
diaBP              0
BMI                19
heartRate          1
glucose            388
TenYearCHD         0
dtype: int64
```

```
In [6]: #managing the null values
mean_value_glucose = df['glucose'].mean()
df['glucose'].fillna(value=mean_value_glucose, inplace=True)
```

```
df['cigsPerDay'].fillna(value=0.0, inplace=True)
np.random.seed(3)
df['education'].fillna(value=np.random.randint(1,5), inplace=True)
mean_value_chol = df['totChol'].mean()
df['totChol'].fillna(value=mean_value_chol, inplace=True)
df['BPMeds'].fillna(value=np.random.randint(0,2), inplace=True)
df['BMI'].fillna(value=np.random.randint(0,2), inplace=True)
mean_value_rate = df['heartRate'].mean()
df['heartRate'].fillna(value=mean_value_rate, inplace=True)
```

```
In [7]: #TenYearCHD means 10-year risk of future coronary heart disease
#TenYearCHD is our target from the training data
y = df['TenYearCHD']
X = df.drop('TenYearCHD', axis = 1)
X = X.to_numpy()
y = y.to_numpy()
```

```
In [8]: def scale_features(X):
min_vec = X.min(axis = 0)
max_vec = X.max(axis = 0)
X = (X - min_vec) / (max_vec - min_vec)
return X, min_vec, max_vec
```

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
print("X_train.shape", X_train.shape, "y_train.shape", y_train.shape)
print("X_test.shape", X_test.shape, "y_test.shape", y_test.shape)
```

```
X_train.shape (3178, 15) y_train.shape (3178,)
X_test.shape (1060, 15) y_test.shape (1060,)
```

```
In [10]: X_train, min_vec, max_vec = scale_features(X_train)
print(X_train)
```

```
[[1.         0.75675676 0.         ... 0.65133929 0.16161616 0.07344633]
 [0.         0.35135135 0.66666667 ... 0.49330357 0.41414141 0.12711864]
 [1.         0.35135135 0.         ... 0.72522321 0.23232323 0.11016949]
 ...
 [1.         0.13513514 0.33333333 ... 0.66629464 0.51515152 0.11855015]
 [0.         0.59459459 0.33333333 ... 0.4921875  0.41414141 0.07627119]
 [0.         0.08108108 0.         ... 0.48058036 0.29292929 0.09887006]]
```

```
In [11]: def sigmoid(z):
g = 1/(1+np.exp(-z))
return g
```

```
In [12]: def compute_cost(X, y, w, b, lambda_ = 1):
m, n = X.shape
cost = 0
for i in range(m):
z = np.dot(X[i],w) + b
f_wb = sigmoid(z)
cost += -y[i]*np.log(f_wb) - (1-y[i])*np.log(1-f_wb)
total_cost = cost/m
return total_cost
np.random.seed(1)
compute_cost(X_train, y_train, 0.01 * np.random.rand(15), 1)
```

Out[12]: 1.1655882597374843

```
In [13]: def compute_gradient(X, y, w, b, lambda_=None):
    m, n = X.shape
    dj_dw = np.zeros(w.shape)
    dj_db = 0.
    for i in range(m):
        f_wb_i = sigmoid(np.dot(X[i],w) + b)
        err_i = f_wb_i - y[i]
        for j in range(n):
            dj_dw[j] = dj_dw[j] + err_i * X[i,j]
        dj_db = dj_db + err_i
    dj_dw = dj_dw/m
    dj_db = dj_db/m
    return dj_db, dj_dw
compute_gradient(X_train, y_train, 0.01 * np.random.rand(15), 1)
```

Out[13]: (0.5771821352171572,
array([0.22677486, 0.24173117, 0.19894651, 0.28046803, 0.06956737,
0.01102355, 0.00188616, 0.15066773, 0.00767253, 0.14903905,
0.16088883, 0.20410731, 0.31590787, 0.18553692, 0.0649162]))

```
In [14]: def gradient_descent(X, y, w_in, b_in, alpha, num_iters, cost_func, gradient_func, l

    for i in range(num_iters) :

        dj_db, dj_dw = gradient_func(X, y, w_in, b_in, lambda_)
        w_in = w_in - alpha * dj_dw
        b_in = b_in - alpha * dj_db

    #         For 10 iterations equally spaced of total iters, print cost.
    if i % math.ceil(num_iters / 10) == 0 or i == (num_iters - 1) :
        print(f"Iteration {i : 4}: Cost {cost_func(X, y, w_in, b_in) : 8.2f}")

    return w_in, b_in
```

```
In [15]: #To verify that gradient descent is working properly, it would be great if we can se
#Running the gradient descent to learn the parameters for our dataset
np.random.seed(1)
initial_w = 0.01* (np.random.rand(15) - 0.5)
print(initial_w)
initial_b = -2
num_iters = 1000
alpha = 0.5
lambda_ = 0
w, b = gradient_descent(X_train, y_train, initial_w, initial_b, alpha, num_iters,com
```

```
[-0.00082978  0.00220324 -0.00499886 -0.00197667 -0.00353244 -0.00407661
 -0.0031374  -0.00154439 -0.00103233  0.00038817 -0.00080805  0.0018522
 -0.00295548  0.00378117 -0.00472612]
Iteration    0: Cost      0.44
Iteration  100: Cost      0.40
Iteration  200: Cost      0.40
Iteration  300: Cost      0.39
Iteration  400: Cost      0.39
Iteration  500: Cost      0.39
Iteration  600: Cost      0.39
Iteration  700: Cost      0.39
Iteration  800: Cost      0.39
```

```
Iteration 900: Cost    0.39
Iteration 999: Cost    0.39
```

```
In [16]: compute_cost(X_train, y_train, w, b)
```

```
Out[16]: 0.3877606709797132
```

```
In [17]: #scaling test data using normalised features
X_test = (X_test - min_vec) / (max_vec - min_vec)
print(X_test)
```

```
[[1.          0.54054054 0.          ... 0.51941964 0.31313131 0.0960452 ]
 [0.          0.2972973  0.          ... 0.49308036 0.26262626 0.09887006]
 [0.          0.2972973  1.          ... 0.36785714 0.31313131 0.0480226 ]
 ...
 [1.          0.94594595 0.          ... 0.63705357 0.41414141 0.1299435 ]
 [0.          0.75675676 0.33333333 ... 0.          0.36363636 0.11855015]
 [0.          0.54054054 0.          ... 0.5171875  0.12121212 0.07909605]]
```

```
In [18]: compute_cost(X_test, y_test, w, b)
```

```
Out[18]: 0.3636064072205817
```

```
In [19]: def predict(X, w, b) :
          m, n = X.shape
          p = np.zeros(m)
          for i in range(m) :
              f_wb = sigmoid(np.dot(X[i], w) + b)
              p[i] = 1 if f_wb > 0.38 else 0
          return p
```

```
In [20]: p = predict(X_train, w,b)
print('Train Accuracy: %f'%(np.mean(p == y_train) * 100))
p = predict(X_test, w,b)
print('Test Accuracy: %f'%(np.mean(p == y_test) * 100))
```

```
Train Accuracy: 83.668974
Test Accuracy: 85.094340
```

```
In [21]: from sklearn.preprocessing import StandardScaler
```

```
In [22]: patient_data = np.array([[1,51,4.0,1,5.0,3.0,0,0,1,400.0,106.0,70.0,26.97,80.0,77.0]
# assume X_train, y_train, X_test, y_test are your training and test data, and w and

# make a prediction for the patient using your trained model
prediction = predict(patient_data, w, b)
print('Prediction: %d'%(prediction))
```

```
Prediction: 1
```